

EmailClassifier

June 25, 2024

0.1 Project: Email Spam Classifier

- Author: Naveed Ur Rehman
- Submitted to : Ezitech

0.2 Overview

This project focuses on building a classifier to distinguish between spam and non-spam emails. Various machine learning models are explored and evaluated for their accuracy in predicting email classifications.

0.2.1 Importing Libraries

```
[ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

## Convert data into vector
from sklearn.feature_extraction.text import CountVectorizer
## Data splitting into train and test set
from sklearn.model_selection import train_test_split

## Dealing with textul data
import nltk
from nltk.corpus import stopwords
from nltk.tokenize import word_tokenize
from nltk.stem import PorterStemmer

# nltk.download('punkt')
# nltk.download('stopwords')

## Models
from sklearn.linear_model import LogisticRegression,SGDClassifier
from sklearn.naive_bayes import GaussianNB,MultinomialNB
from sklearn.ensemble import RandomForestClassifier
```

```
from sklearn.metrics import
    accuracy_score, classification_report, confusion_matrix
```

0.2.2 Importing datasets

```
[ ]: df = pd.read_csv("data/emails.csv")
df.head()
```

```
[ ]:
      text  spam
0  Subject: naturally irresistible your corporate...    1
1  Subject: the stock trading gunslinger fanny i...    1
2  Subject: unbelievable new homes made easy im ...    1
3  Subject: 4 color printing special request add...    1
4  Subject: do not have money , get software cds ...    1
```

```
[ ]: df1 = pd.read_csv("./data/spam_ham_dataset.csv")
df1.head()
```

```
[ ]:
      Unnamed: 0  label      text \
0          605    ham  Subject: enron methanol ; meter # : 988291\r\n...
1         2349    ham  Subject: hpl nom for january 9 , 2001\r\n( see...
2         3624    ham  Subject: neon retreat\r\nho ho ho , we ' re ar...
3         4685  spam  Subject: photoshop , windows , office . cheap ...
4          2030    ham  Subject: re : indian springs\r\nthis deal is t...

      label_num
0              0
1              0
2              0
3              1
4              0
```

```
[ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5728 entries, 0 to 5727
Data columns (total 2 columns):
#   Column  Non-Null Count  Dtype
---  -
0    text    5728 non-null    object
1    spam    5728 non-null    int64
dtypes: int64(1), object(1)
memory usage: 89.6+ KB
```

```
[ ]: df1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```

RangeIndex: 5171 entries, 0 to 5170
Data columns (total 4 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Unnamed: 0   5171 non-null   int64
1   label        5171 non-null   object
2   text         5171 non-null   object
3   label_num    5171 non-null   int64
dtypes: int64(2), object(2)
memory usage: 161.7+ KB

```

```
[ ]: df.shape,df1.shape
```

```
[ ]: ((5728, 2), (5171, 4))
```

0.2.3 Data cleaning

```
[ ]: df1.drop(['Unnamed: 0','label'],axis=1,inplace=True)
df1
```

```
[ ]:
      text  label_num
0  Subject: enron methanol ; meter # : 988291\r\n...      0
1  Subject: hpl nom for january 9 , 2001\r\n( see...      0
2  Subject: neon retreat\r\nho ho ho , we ' re ar...      0
3  Subject: photoshop , windows , office . cheap ...      1
4  Subject: re : indian springs\r\nthis deal is t...      0
...
5166 Subject: put the 10 on the ft\r\nthe transport...      0
5167 Subject: 3 / 4 / 2000 and following noms\r\nhnp...      0
5168 Subject: calpine daily gas nomination\r\n>\r\n...      0
5169 Subject: industrial worksheets for august 2000...      0
5170 Subject: important online banking alert\r\nndea...      1

[5171 rows x 2 columns]
```

```
[ ]: df1.rename({"label_num":"spam"},axis=1,inplace=True)
```

```
[ ]: sum(df.duplicated()),sum(df1.duplicated())
```

```
[ ]: (33, 178)
```

```
[ ]: df.shape,df1.shape
```

```
[ ]: ((5728, 2), (5171, 2))
```

```
[ ]: df.drop_duplicates(inplace=True)
df1.drop_duplicates(inplace=True)
```

```
[ ]: df.shape, df1.shape
```

```
[ ]: ((5695, 2), (4993, 2))
```

0.2.4 Data Merging

```
[ ]: combine_df = pd.concat([df, df1])  
      combine_df.head()
```

```
[ ]: 

|   | text                                              | spam |
|---|---------------------------------------------------|------|
| 0 | Subject: naturally irresistible your corporate... | 1    |
| 1 | Subject: the stock trading gunslinger fanny i...  | 1    |
| 2 | Subject: unbelievable new homes made easy im ...  | 1    |
| 3 | Subject: 4 color printing special request add...  | 1    |
| 4 | Subject: do not have money , get software cds ... | 1    |



```

```
[ ]: sum(combine_df.duplicated())
```


```
[ ]: 0
```

0.2.5 Data Visualization

```
[ ]: plt.hist(combine_df['spam'])  
      plt.title("Spam vs Not Spam")  
      plt.show()
```



```
[ ]: df['spam'].plot.line()  
plt.title("Spam vs Not Spam")  
plt.show()
```



0.2.6 Preprocessing

- Convert to lowercase
- Remove Unnecessary characters such as Punctuation and Special Characters
- Remove stop words
- Perform stemming or lemmatization
- Handle missing values if any
- There is no missing data point in the dataset

```
[ ]: # Remove Unnecessary characters such as Punctuation and Special Characters
import re
import string

def remove_punc_spec_chars(text):
    # Remove punctuation
    text = text.translate(str.maketrans("", "", string.punctuation))

    # Remove special characters (keeping alphanumeric characters and spaces)
    text = re.sub(r"[^a-zA-Z0-9\s]", "", text)

    # Remove extra whitespace
```

```
text = re.sub("\s+", " ",text).strip()
return text
```

```
[ ]: def preprocess_text(text):
    # Convert to lowercase
    text = text.lower()

    # Remove punctuation and special characters
    text = remove_punc_spec_chars(text)

    # Tokenize
    tokens = word_tokenize(text)

    # Remove stopwords
    stop_words = set(stopwords.words('english'))
    tokens = [word for word in tokens if word not in stop_words]

    # Stemming
    stemmer = PorterStemmer()
    tokens = [stemmer.stem(word) for word in tokens]

    # Join tokens back into a string
    preprocessed_text = ' '.join(tokens)

    return preprocessed_text
```

```
[ ]: combine_df['text'] = combine_df['text'].apply(preprocess_text)
```

```
[ ]: combine_df
```

```
[ ]:
```

	text	spam
0	subject natur irresist corpor ident lt realli ...	1
1	subject stock trade gunsling fanni merril muzo...	1
2	subject unbeliev new home made easi im want sh...	1
3	subject 4 color print special request addit in...	1
4	subject money get softwar cd softwar compat gr...	1
...
5165	subject fw crosstex energi driscoll ranch 1 3 m...	0
5166	subject put 10 ft transport volum decreas 2500...	0
5167	subject 3 4 2000 follow nom hpl take extra 15 ...	0
5169	subject industri worksheet august 2000 activ a...	0
5170	subject import onlin bank alert dear valu citi...	1

[10688 rows x 2 columns]

0.2.7 Split the data into Train and Test set

```
[ ]: X = combine_df['text']
     y = combine_df['spam']

[ ]: # Split the data
     X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
     ↪random_state=42)
```

0.2.8 Data Encoding

```
[ ]: # Convert the data into vectors
     vectorizer = CountVectorizer()
     X_train_transform = vectorizer.fit_transform(X_train)
     X_test_transform = vectorizer.transform(X_test)

[ ]: X_train_transform = X_train_transform.toarray()

[ ]: X_train_transform.shape, y_train.shape,

[ ]: ((7481, 49571), (7481,))
```

0.2.9 Training Model

```
[ ]: clf1 = LogisticRegression(max_iter=1000)
     clf2 = RandomForestClassifier()
     clf3 = SGDClassifier(max_iter=1000, loss='hinge')
     clf4 = GaussianNB()
     clf5 = MultinomialNB()

[ ]: clf1.fit(X_train_transform, y_train)

[ ]: LogisticRegression(max_iter=1000)

[ ]: clf2.fit(X_train_transform, y_train)

[ ]: RandomForestClassifier()

[ ]: clf3.fit(X_train_transform, y_train)

[ ]: SGDClassifier()

[ ]: clf4.fit(X_train_transform, y_train)
     clf5.fit(X_train_transform, y_train)

[ ]: MultinomialNB()
```


0.2.10 Model Prediction on test data

```
[ ]: clf1_pred = clf1.predict(X_test_transform)
      clf2_pred = clf2.predict(X_test_transform)
      clf3_pred = clf3.predict(X_test_transform)
```

```
[ ]: clf4_pred = clf4.predict(X_test_transform.toarray())
```

```
[ ]: clf5_pred = clf5.predict(X_test_transform)
```

0.2.11 Model Evaluation

```
[ ]: print("Accuracy Score Clf1:",accuracy_score(clf1_pred,y_test))
      print("Accuracy Score Clf4:",accuracy_score(clf2_pred,y_test))
      print("Accuracy Score Clf6:",accuracy_score(clf3_pred,y_test))
      print("Accuracy Score Clf6:",accuracy_score(clf4_pred,y_test))
      print("Accuracy Score Clf6:",accuracy_score(clf5_pred,y_test))
```

```
Accuracy Score Clf1: 0.9856563766760212
Accuracy Score Clf4: 0.9766136576239476
Accuracy Score Clf6: 0.9787963829123791
Accuracy Score Clf6: 0.9513564078578111
Accuracy Score Clf6: 0.9840972871842844
```

```
[ ]: print("Accuracy Score Clf1:",confusion_matrix(clf1_pred,y_test))
      print("Accuracy Score Clf4:",confusion_matrix(clf2_pred,y_test))
      print("Accuracy Score Clf6:",confusion_matrix(clf3_pred,y_test))
      print("Accuracy Score Clf6:",confusion_matrix(clf4_pred,y_test))
      print("Accuracy Score Clf6:",confusion_matrix(clf5_pred,y_test))
```

```
Accuracy Score Clf1: [[2351   21]
 [  25  810]]
Accuracy Score Clf4: [[2342   41]
 [  34  790]]
Accuracy Score Clf6: [[2346   38]
 [  30  793]]
Accuracy Score Clf6: [[2349  129]
 [  27  702]]
Accuracy Score Clf6: [[2350   25]
 [  26  806]]
```

```
[ ]: print("Accuracy Score Clf1:",classification_report(clf1_pred,y_test))
      print("Accuracy Score Clf4:",classification_report(clf2_pred,y_test))
      print("Accuracy Score Clf6:",classification_report(clf3_pred,y_test))
      print("Accuracy Score Clf6:",classification_report(clf4_pred,y_test))
      print("Accuracy Score Clf6:",classification_report(clf5_pred,y_test))
```

```
Accuracy Score Clf1:          precision    recall  f1-score   support
```

0	0.99	0.99	0.99	2372		
1	0.97	0.97	0.97	835		
accuracy			0.99	3207		
macro avg	0.98	0.98	0.98	3207		
weighted avg	0.99	0.99	0.99	3207		

Accuracy Score Clf4:			precision	recall	f1-score	support
0	0.99	0.98	0.98	2383		
1	0.95	0.96	0.95	824		
accuracy			0.98	3207		
macro avg	0.97	0.97	0.97	3207		
weighted avg	0.98	0.98	0.98	3207		

Accuracy Score Clf6:			precision	recall	f1-score	support
0	0.99	0.98	0.99	2384		
1	0.95	0.96	0.96	823		
accuracy			0.98	3207		
macro avg	0.97	0.97	0.97	3207		
weighted avg	0.98	0.98	0.98	3207		

Accuracy Score Clf6:			precision	recall	f1-score	support
0	0.99	0.95	0.97	2478		
1	0.84	0.96	0.90	729		
accuracy			0.95	3207		
macro avg	0.92	0.96	0.93	3207		
weighted avg	0.96	0.95	0.95	3207		

Accuracy Score Clf6:			precision	recall	f1-score	support
0	0.99	0.99	0.99	2375		
1	0.97	0.97	0.97	832		
accuracy			0.98	3207		
macro avg	0.98	0.98	0.98	3207		
weighted avg	0.98	0.98	0.98	3207		

0.2.12 Visualizing Accuracies of Models

```
[ ]: import matplotlib.pyplot as plt
import numpy as np

def plot_model_accuracies(models, scores):
    # Ensure models and scores are lists
    if not isinstance(models, list) or not isinstance(scores, list):
        raise ValueError("Both models and scores should be lists")

    # Ensure the lengths match
    if len(models) != len(scores):
        raise ValueError("The number of models and scores should be the same")

    # Create the plot
    fig, ax = plt.subplots(figsize=(10, 6))

    # Create the bars
    bars = ax.bar(models, scores, color='skyblue', edgecolor='navy')

    # Customize the plot
    ax.set_ylabel('Accuracy Score')
    ax.set_title('Comparison of Model Accuracies')
    ax.set_ylim(0, 1) # Assuming accuracy scores are between 0 and 1

    # Add value labels on the bars
    for bar in bars:
        height = bar.get_height()
        ax.text(bar.get_x() + bar.get_width()/2., height,
                f'{height:.3f}',
                ha='center', va='bottom')

    # Add a horizontal line for the mean accuracy
    mean_accuracy = np.mean(scores)
    ax.axhline(y=mean_accuracy, color='red', linestyle='--', label=f'Mean_
↳Accuracy: {mean_accuracy:.3f}')

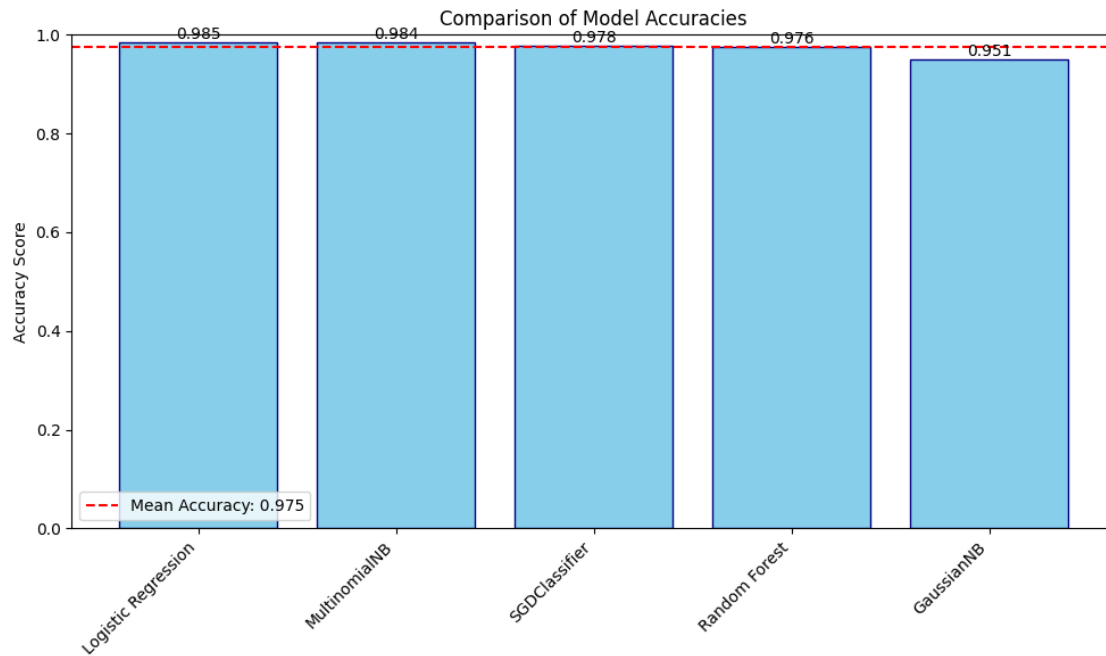
    # Rotate x-axis labels for better readability
    plt.xticks(rotation=45, ha='right')

    # Add legend
    plt.legend()

    # Adjust layout and display the plot
    plt.tight_layout()
    plt.show()
```

```
# Example usage:
models = ['Logistic Regression', 'MultinomialNB', 'SGDClassifier', 'Random_
↳Forest', 'GaussianNB']
accuracy_scores = [0.985, 0.984, 0.978, 0.976, 0.951] # Replace with your_
↳actual accuracy scores

plot_model_accuracies(models, accuracy_scores)
```



0.2.13 Save the trained models into file

```
[ ]: import pickle
models = [clf1, clf2, clf3, clf4, clf5]
for model in models:
    with open(f"{model}.pkl", 'wb') as file:
        pickle.dump(model, file)
```

0.2.14 Conclusion:

The Logistic Regression model demonstrated the highest accuracy of 98.5% in classifying emails as spam or non-spam. This project effectively illustrates the process of building and evaluating a text classification model for email spam detection.